

Hands-on Lab: Populating a Data Warehouse using PostgreSQL



Estimated time needed: 1 hour

Purpose of the Lab:

The lab is designed to provide hands-on experience in creating and managing a production database using PostgreSQL within the IBM Skills Network Labs (SN Labs) Cloud IDE. You will learn how to launch a PostgreSQL server instance, utilize the pgAdmin graphical user interface (GUI) for database operations, and execute essential tasks like creating a database, designing tables, and loading data. The lab focuses on building a foundation in database management by guiding learners through the process of setting up a 'Production' database and populating it with data following a star schema design.

Benefits of Learning the Lab:

Engaging in this lab offers significant benefits for learners seeking to deepen their understanding of database management systems, particularly PostgreSQL. By working through the lab, you will gain practical skills in SQL, database creation, table design, and data manipulation, which are crucial for roles in data engineering, database administration, and data science. The hands-on approach helps in consolidating knowledge of database schemas and SQL queries, thereby enhancing the learner's ability to manage and analyze data effectively in real-world scenarios. Additionally, familiarity with tools like pgAdmin and the Cloud IDE environment adds valuable experience to your skill set, preparing you for advanced database projects and tasks.

Software Used in this Lab

To complete this lab you will utilize the [PostgreSQL Database](#) relational database service available as part of IBM Skills Network Labs (SN Labs) Cloud IDE. SN Labs is a virtual lab environment used in this course.



Database Used in this Lab

Production database is used in this lab.
The production database contains:

- DimCustomer
- DimMonth
- FactBilling

Objectives

In this lab you will:

- Create production related database and tables in a PostgreSQL instance.
- Populate the production data warehouse byloading the tables from Scripts.

Lab Structure

In this lab, you will complete several tasks in which you will learn how to create tables and load data in the PostgreSQL database service using the pgAdmin graphical user interface (GUI) tool.

Data Used in this Lab

The following are the SQL data files used in this lab.

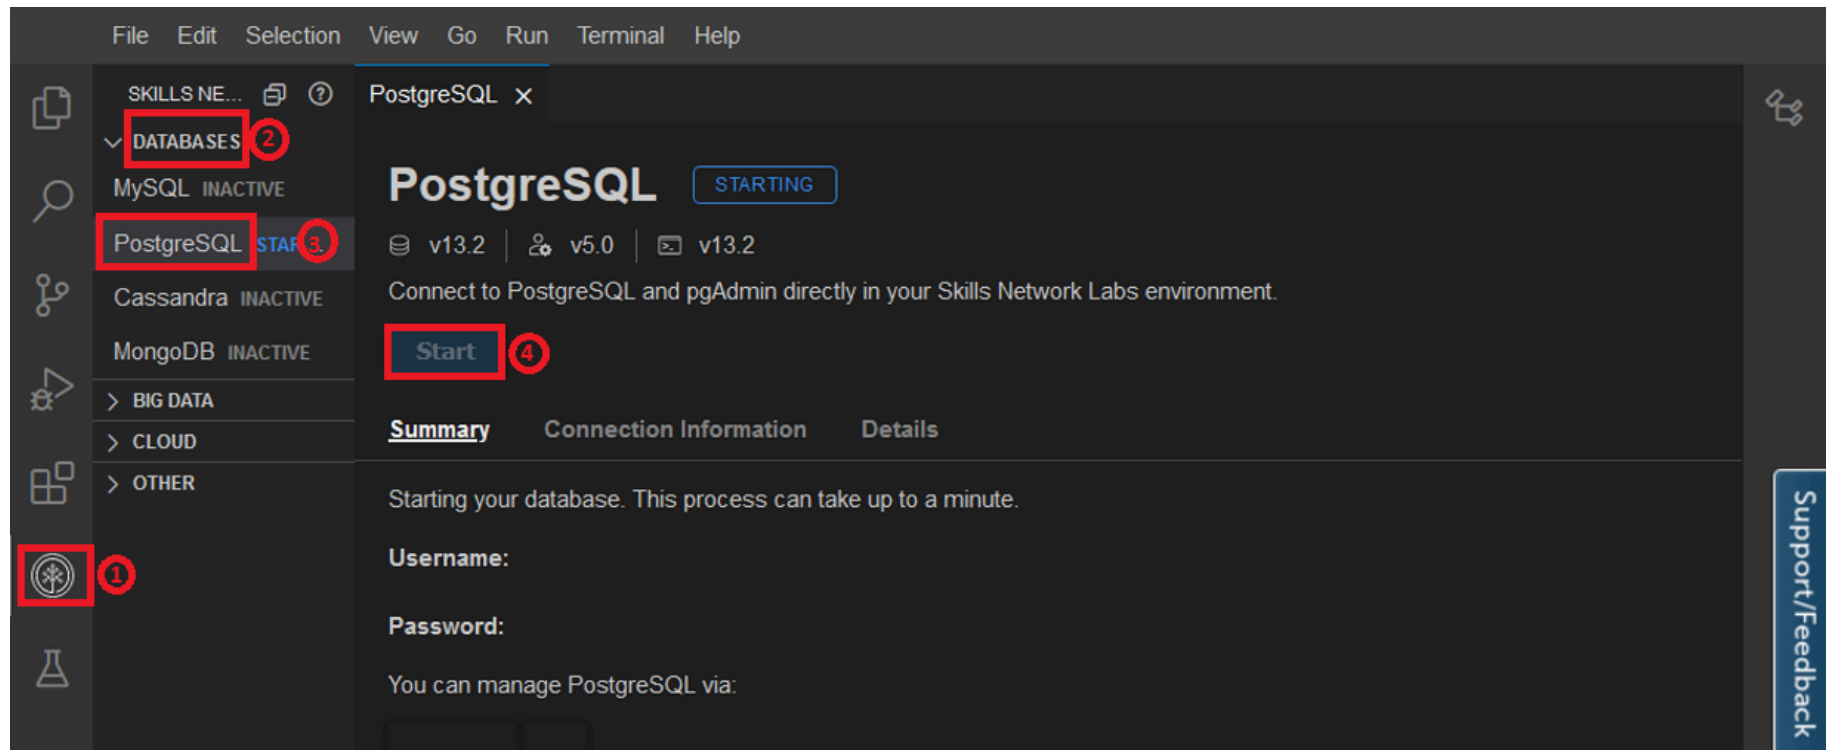
The production database contains:

- [DimCustomer](#)
- [DimMonth](#)
- [FactBilling](#)
- [Star Schema](#)

Task A: Create a database

First, to create a database on a PostgreSQL server instance, you'll first want to actually launch a PostgreSQL server instance on Cloud IDE and open up the pgAdmin Graphical User Interface.

1. Click on the Skills Network extension button on the left side of the window.
2. Open the **DATABASES** drop down menu.
3. Click on **PostgreSQL**
4. Click on the **Start** button. PostgreSQL may take a few moments to start.



5. Next, open the pgAdmin Graphical User Interface by clicking the **pgAdmin** button in the Cloud IDE interface.

The screenshot shows the IBM Skills Network Labs interface for PostgreSQL. The top bar includes 'Lab', 'IBMCLOUD', and 'Launch Application'. Below this is a menu with 'File', 'Edit', 'Selection', 'View', 'Go', 'Run', 'Terminal', and 'Help'. On the left, a sidebar lists various databases: MySQL (INACTIVE), PostgreSQL (ACTIVE), Cassandra (INACTIVE), and MongoDB (INACTIVE). The main panel displays 'PostgreSQL' with an 'ACTIVE' status. It shows version information (v13.2, v5.0, v13.2) and a 'Stop' button. Below this, there are tabs for 'Summary', 'Connection Information', and 'Details'. The 'Summary' tab is selected, showing instructions on how to connect to PostgreSQL and pgAdmin. It includes fields for 'Username' and 'Password', both with copy icons. Below these fields, it says 'You can manage PostgreSQL via:' and highlights the 'pgAdmin' button with a red box. At the bottom, there are buttons for 'PostgreSQL CLI' and 'New Terminal'. A 'Support/Feedback' button is located on the right side of the interface.

Lab IBMCLOUD Launch Application

File Edit Selection View Go Run Terminal Help

SKILLS NE... ? PostgreSQL x

✓ DATABASES

MySQL INACTIVE

PostgreSQL ACTIVE

Cassandra INACTIVE

MongoDB INACTIVE

> BIG DATA

> CLOUD

> OTHER

PostgreSQL

ACTIVE

v13.2 | v5.0 | v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary Connection Information Details

Your database and pgAdmin server are now ready to use and available with the following login credentials. For more details on how to navigate PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

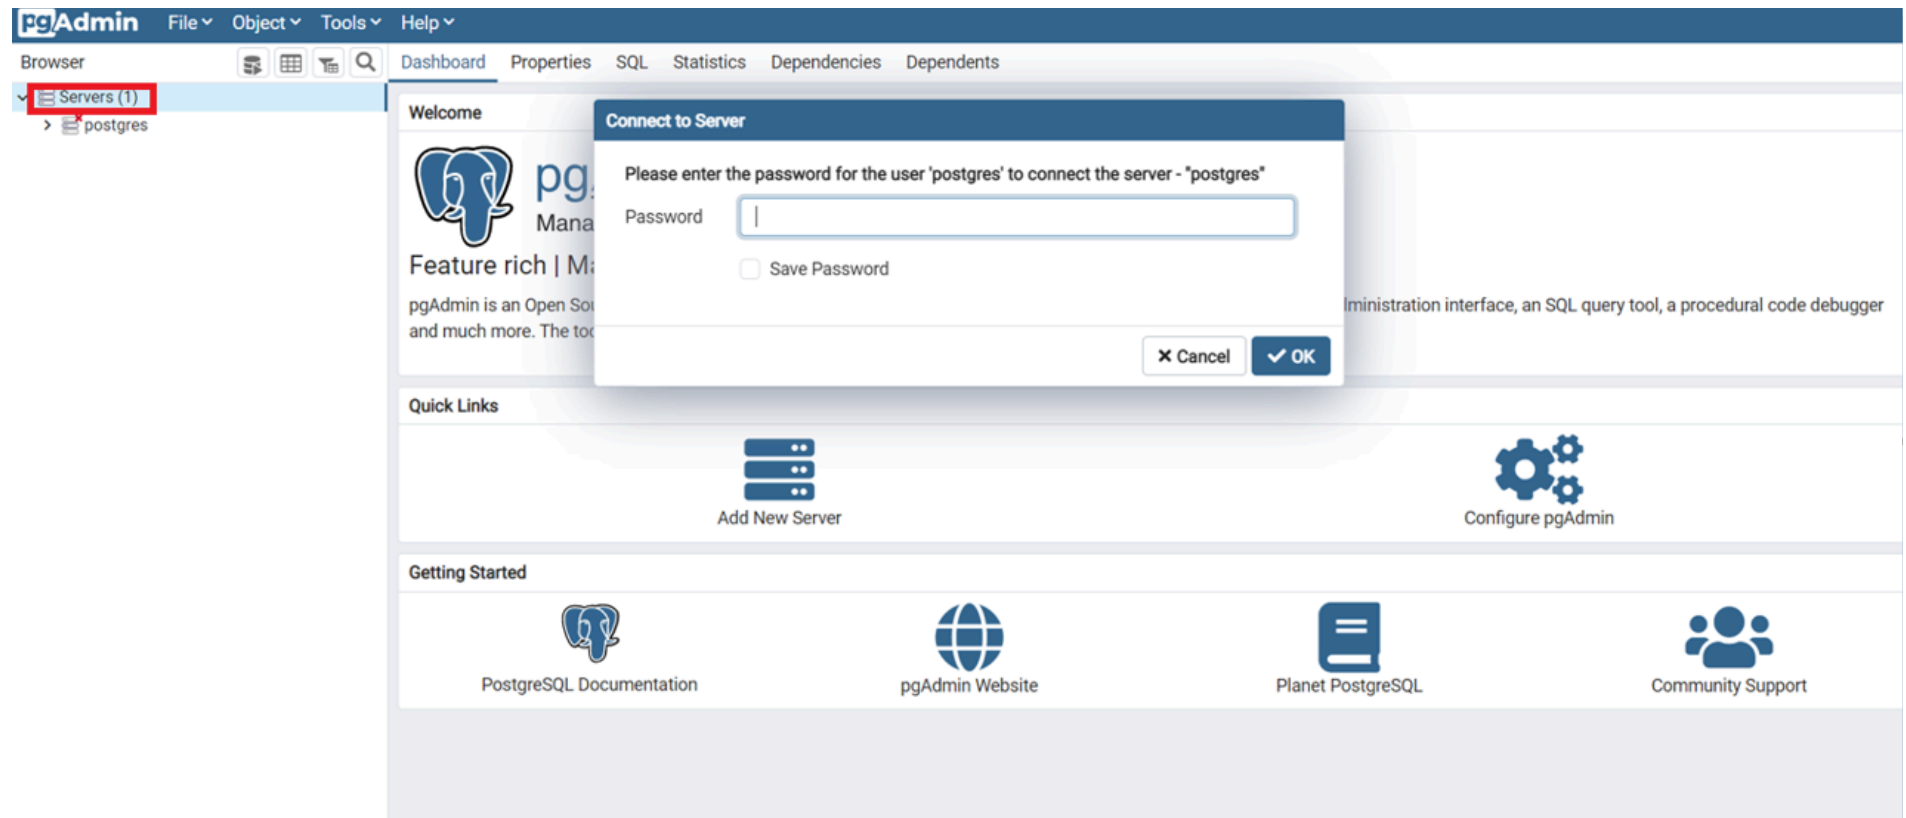
pgAdmin

Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI New Terminal

Support/Feedback

6. Once the pgAdmin GUI opens, click on the **Servers** tab on the left side of the page. You will be prompted to enter a password.



7. To retrieve your password, click on the **PostgreSQL** tab near the top of the interface.
8. Click on the **Copy** icon to the right of your password to copy the session password onto your clipboard.

The screenshot shows the IBM Skills Network Labs interface for PostgreSQL. The top navigation bar includes 'Lab', 'IBMCLOUD', and 'Launch Application'. Below this is a menu bar with 'File', 'Edit', 'Selection', 'View', 'Go', 'Run', 'Terminal', and 'Help'. The left sidebar contains a list of databases: MySQL (INACTIVE), PostgreSQL (ACTIVE), Cassandra (INACTIVE), and MongoDB (INACTIVE). The main content area is titled 'PostgreSQL' and shows it is 'ACTIVE'. It includes version information (v13.2, v5.0, v13.2) and a 'Stop' button. Below this are tabs for 'Summary', 'Connection Information', and 'Details'. The 'Summary' tab is selected, showing instructions on how to connect to PostgreSQL and pgAdmin. It displays the 'Username' and 'Password' fields, with the password field highlighted by a red square. Below the password field, there is a link to 'pgAdmin' and a button to 'New Terminal'.

Lab IBMCLOUD Launch Application

File Edit Selection View Go Run Terminal Help

SKILLS NE... PostgreSQL x

DATABASES

MySQL INACTIVE

PostgreSQL ACTIVE

Cassandra INACTIVE

MongoDB INACTIVE

BIG DATA

CLOUD

OTHER

PostgreSQL

ACTIVE

v13.2 v5.0 v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary Connection Information Details

Your database and pgAdmin server are now ready to use and available with the following login credentials. For more details on how to navigate PostgreSQL, please check out the Details section.

Username: [REDACTED]

Password: [REDACTED]

You can manage PostgreSQL via:

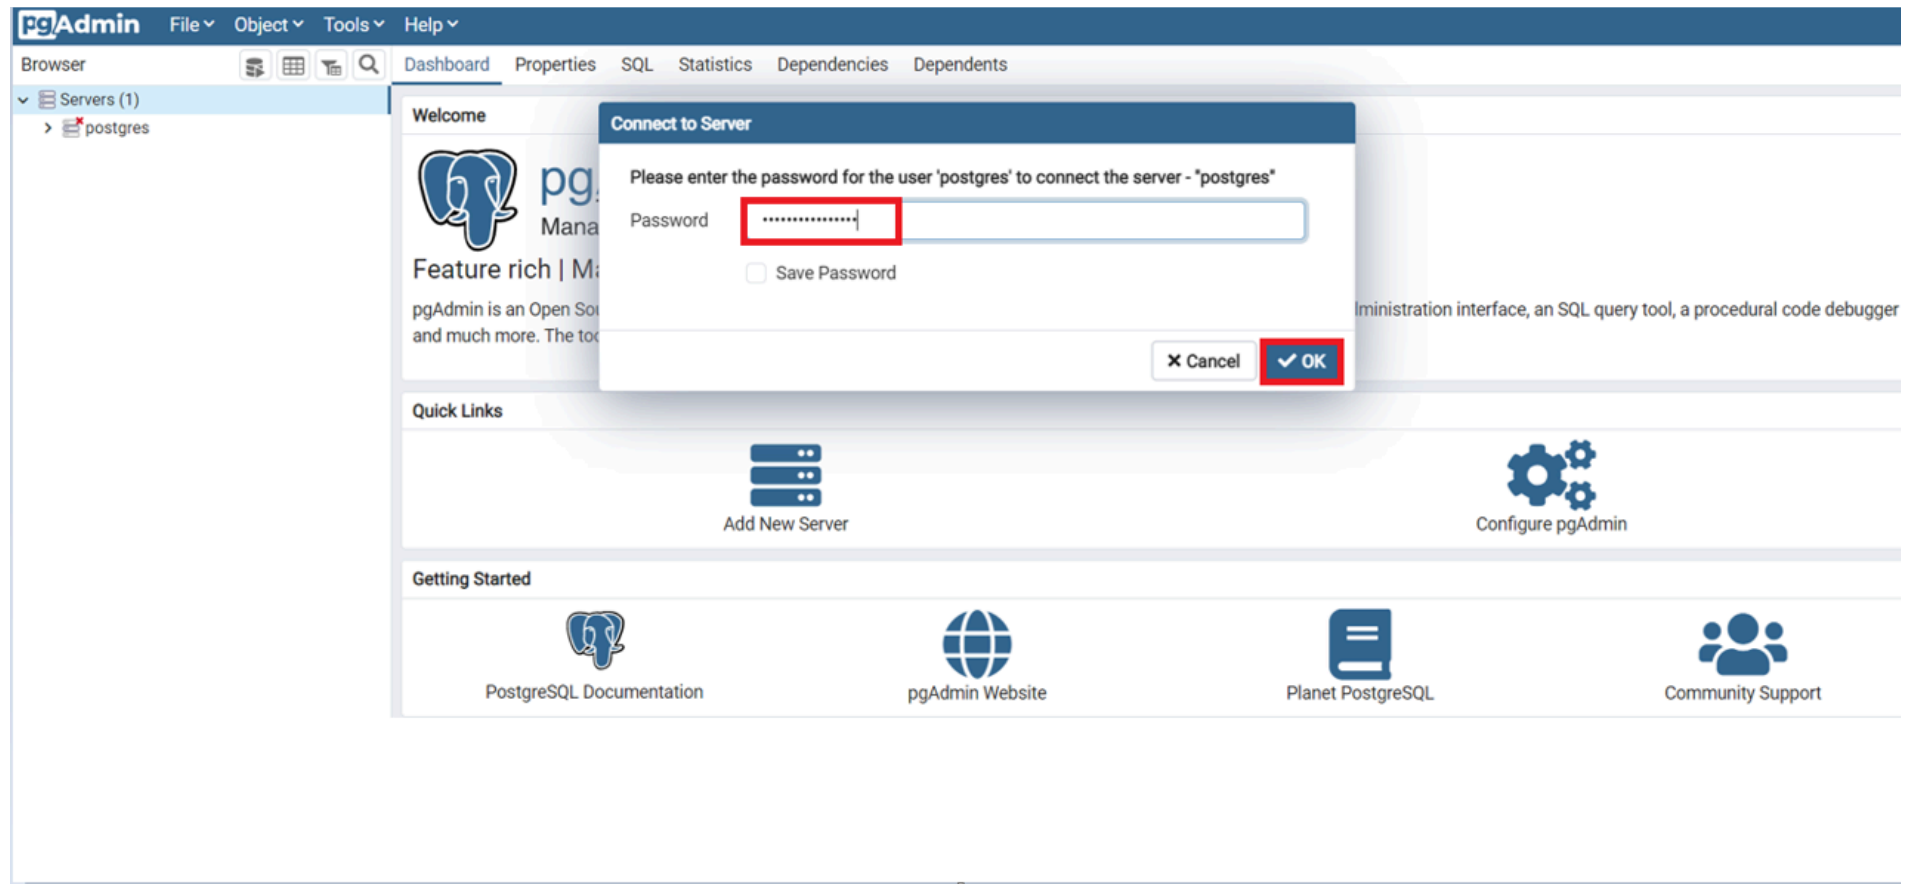
pgAdmin

Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI New Terminal

Support/Feedback

9. Navigate back to the **pgAdmin** tab and paste in your password, then click OK.



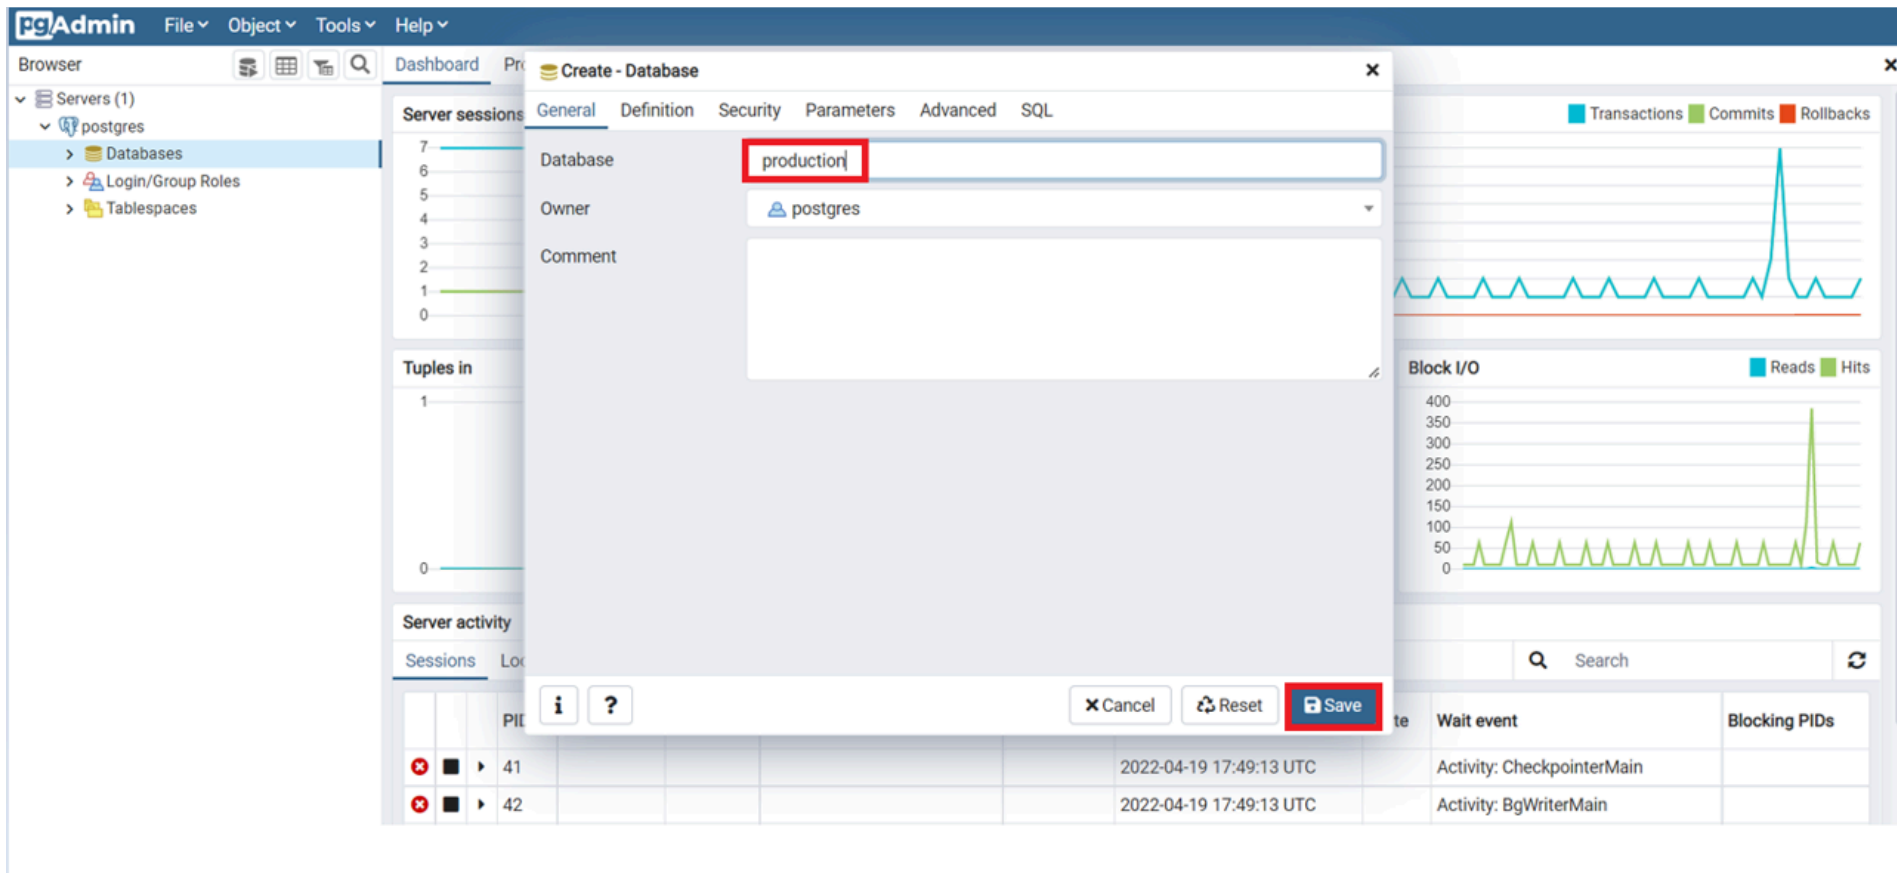
10. You will then be able to access the pgAdmin GUI tool.

11. In the left tree-view, right-click on **Databases**> **Create** > **Database**.

The screenshot shows the pgAdmin 4 interface. In the left sidebar, the 'Databases' folder is selected. A context menu is open, showing the 'Create' button and the 'Database...' option. The 'Create' button is highlighted with a red box and a red circle. The 'Database...' option is also highlighted with a red box and a red circle. The 'Server sessions' chart shows 7 total sessions, with 1 active and 6 idle. The 'Transactions per second' chart shows a peak of 3 transactions per second. The 'Tuples in' chart shows 1 insert. The 'Tuples out' chart shows 1400 fetched tuples. The 'Block I/O' chart shows 350 reads. The 'Server activity' table at the bottom shows three active sessions.

Server activity										
Sessions										
		PID	Database	User	Application	Client	Backend start	State	Wait event	Blocking PIDs
✖	■	▶ 41					2022-04-19 17:49:13 UTC		Activity: CheckpointerMain	
✖	■	▶ 42					2022-04-19 17:49:13 UTC		Activity: BgWriterMain	
✖	■	▶ 43					2022-04-19 17:49:13 UTC		Activity: WalWriterMain	

In the Database box, type **Production** as the name for your new database, and then click **Save**. Proceed to Task B.

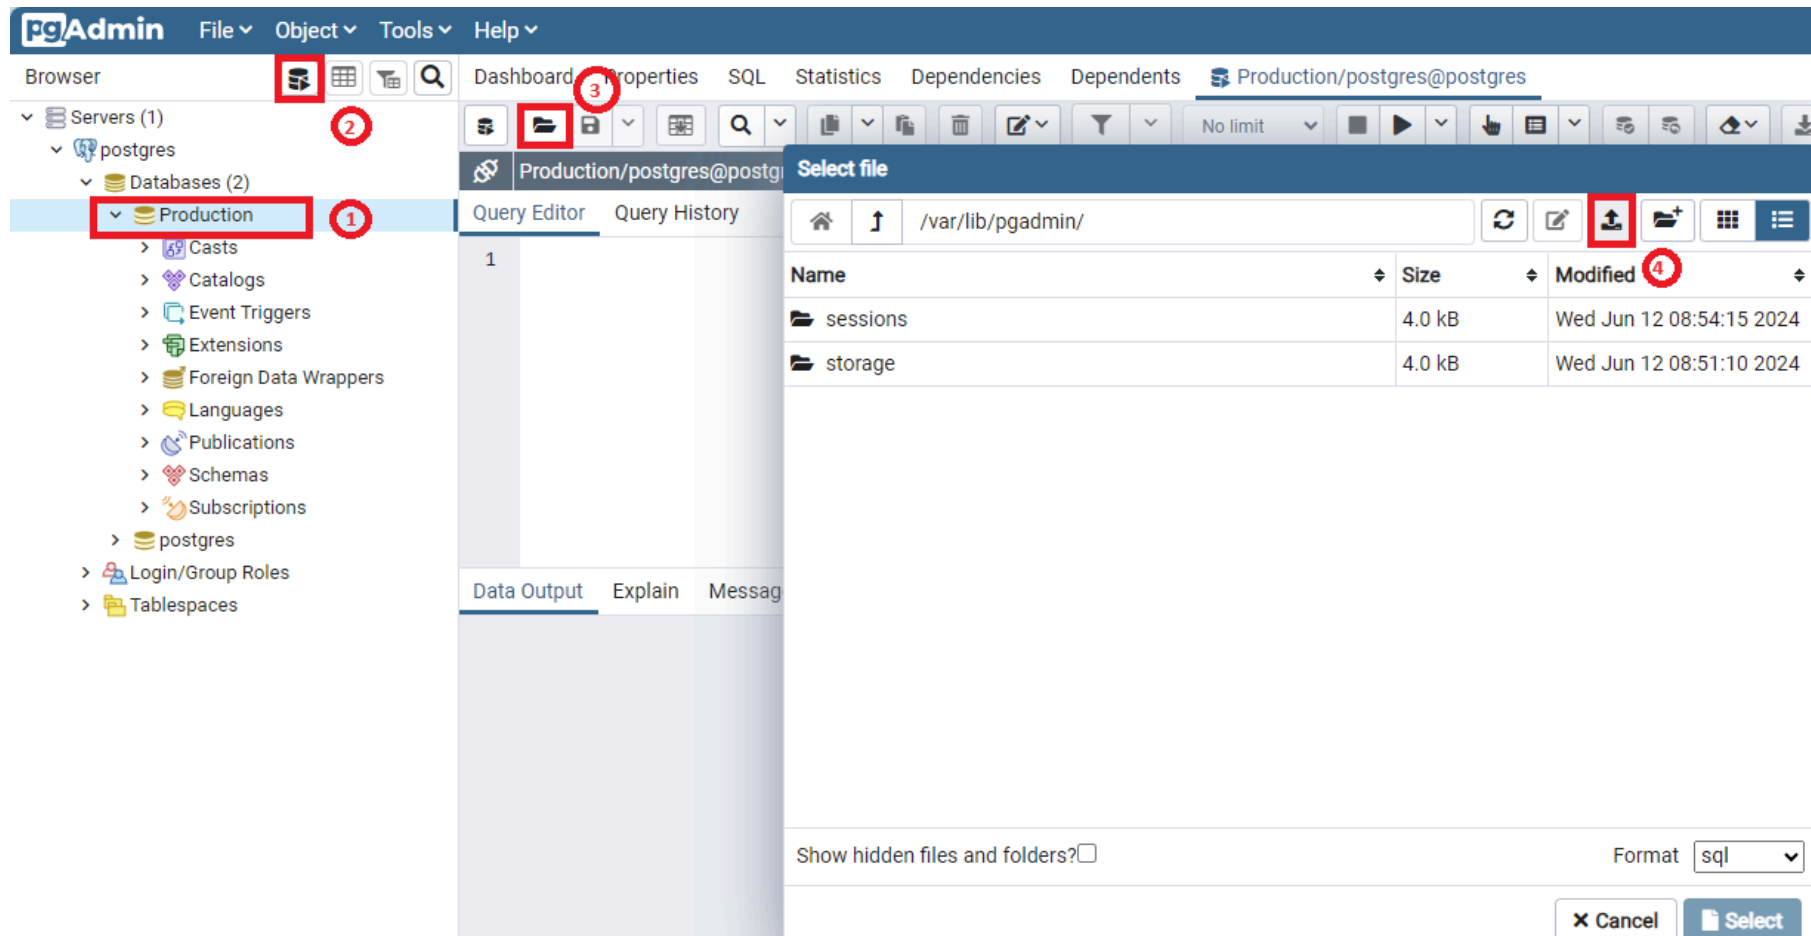


Task B: Create tables

Now, that you have your PostgreSQL service active and have created the **Production database** using pgAdmin, let's go ahead and create a few tables to populate the database and store the data that we wish to eventually upload into it.

1. Click on the Production database and in the top of the page go to **Query tool** and then click on **Open File**. Next a new page pops up called **Select File**. Click on **Upload** icon as shown in the screenshot.

Note: Ensure that you upload the files to this path: /var/lib/pgadmin/



The screenshot shows the pgAdmin interface with the following elements:

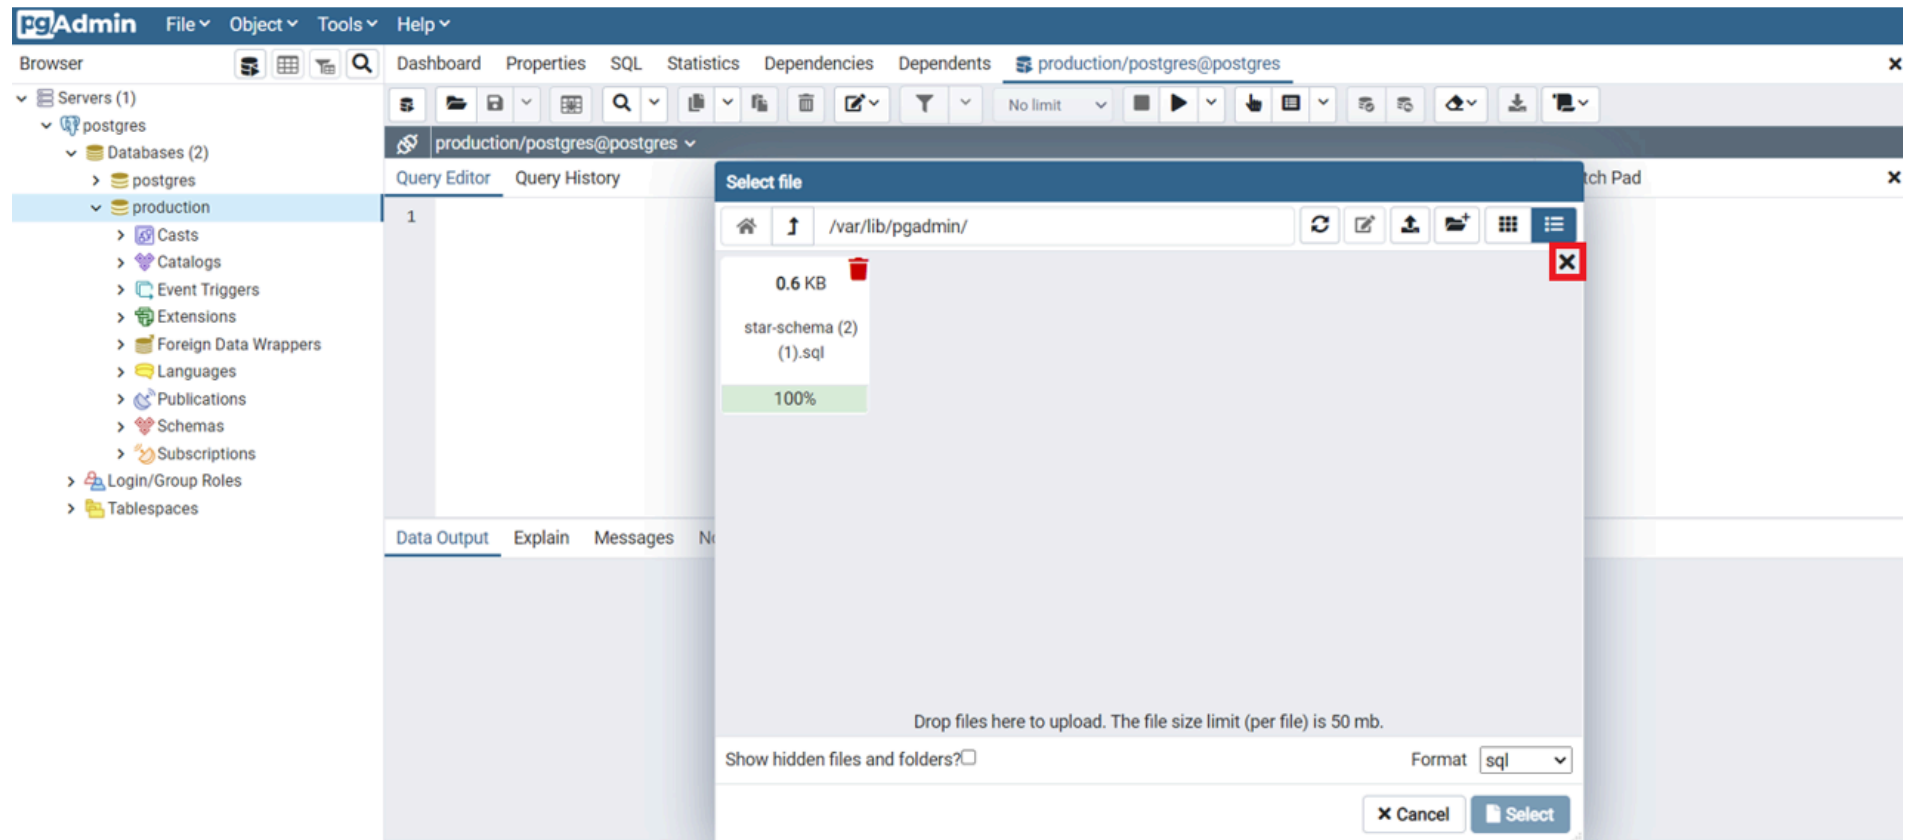
- Left Sidebar:** The 'Production' database is selected under 'Databases (2)' (1).
- File List:** The 'star-schema.sql' file is highlighted in the list (2).
- Upload Button:** The 'Upload' button (represented by an upward arrow) is clicked (3).
- Confirmation:** The 'X' icon on the right side of the page is clicked to confirm the upload (4).

The 'Select file' dialog shows the following table:

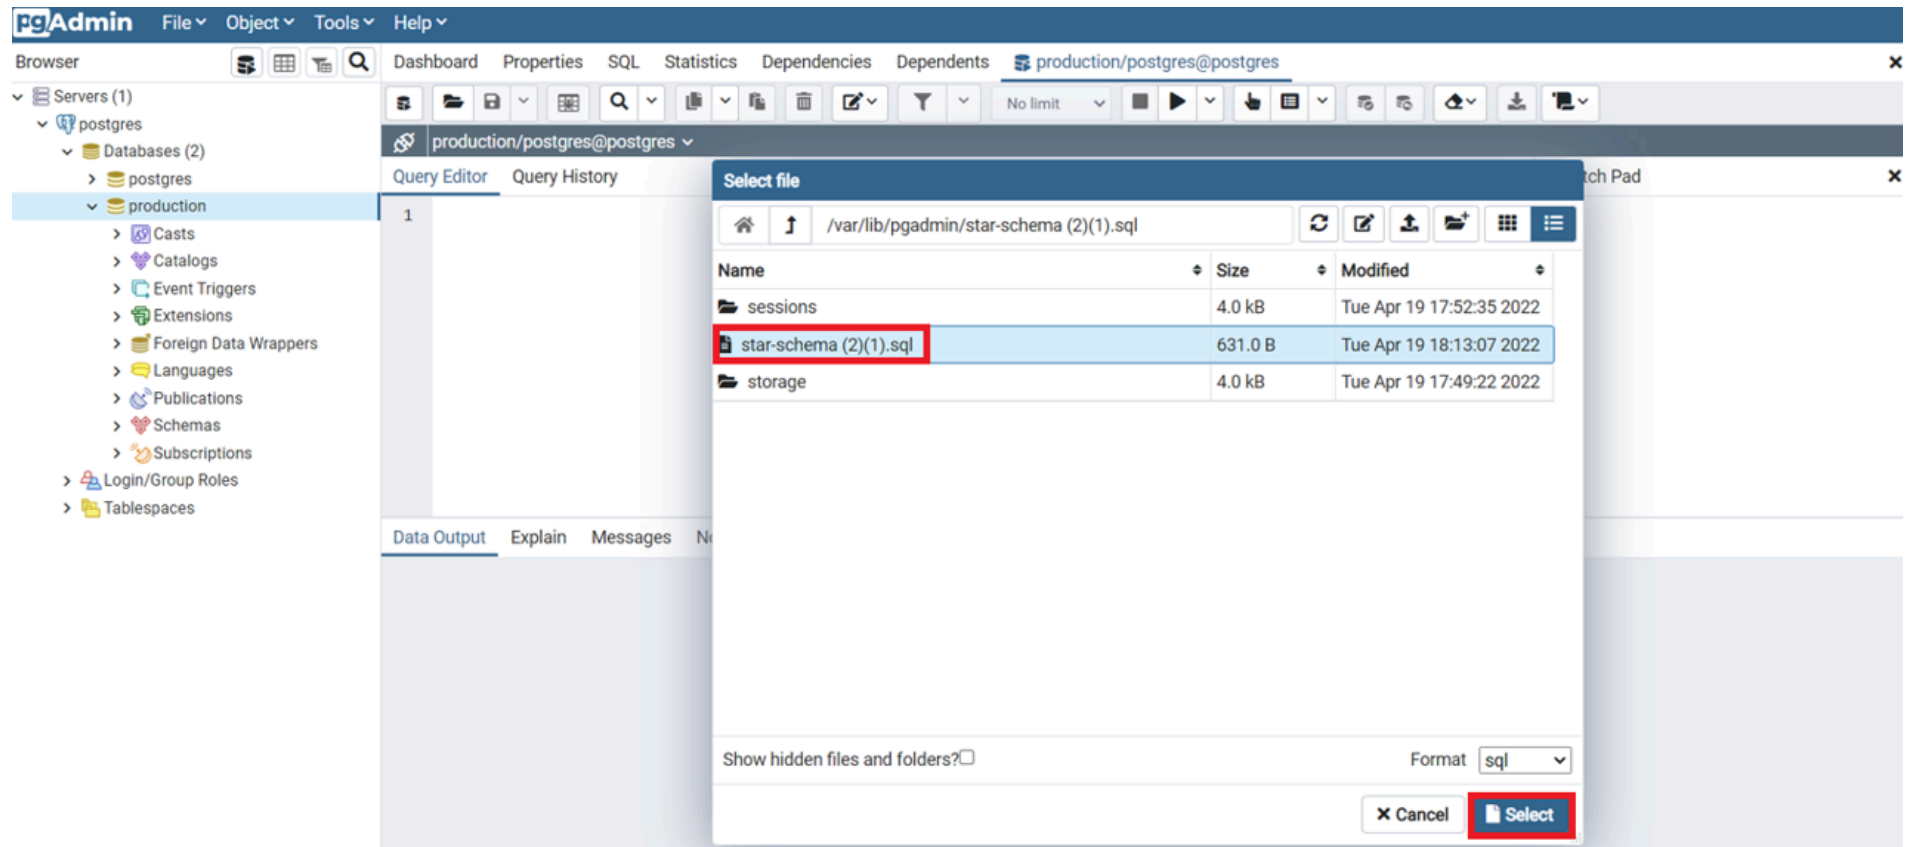
Name	Size	Modified
sessions	4.0 kB	Wed Jun 12 08:54:15 2024
storage	4.0 kB	Wed Jun 12 08:51:10 2024

At the bottom, there is a checkbox for 'Show hidden files and folders?' and a 'Format' dropdown set to 'sql'. Buttons for 'Cancel' and 'Select' are also present.

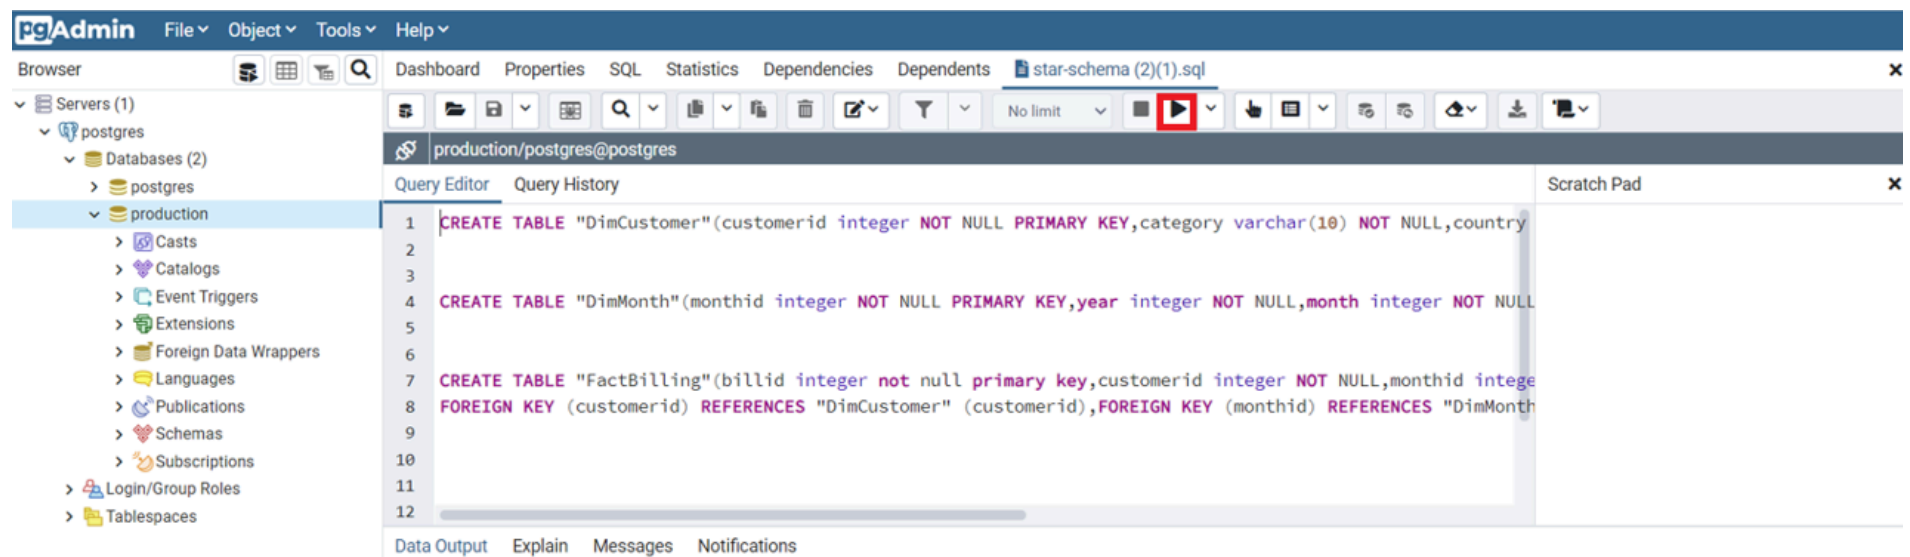
2. In the new blank page that appears drag and drop the **star-schema.sql** file inside the blank page. Once the **star-schema.sql** file is successfully loaded, click on the **X** icon on the right hand side of the page as shown in the screenshot.



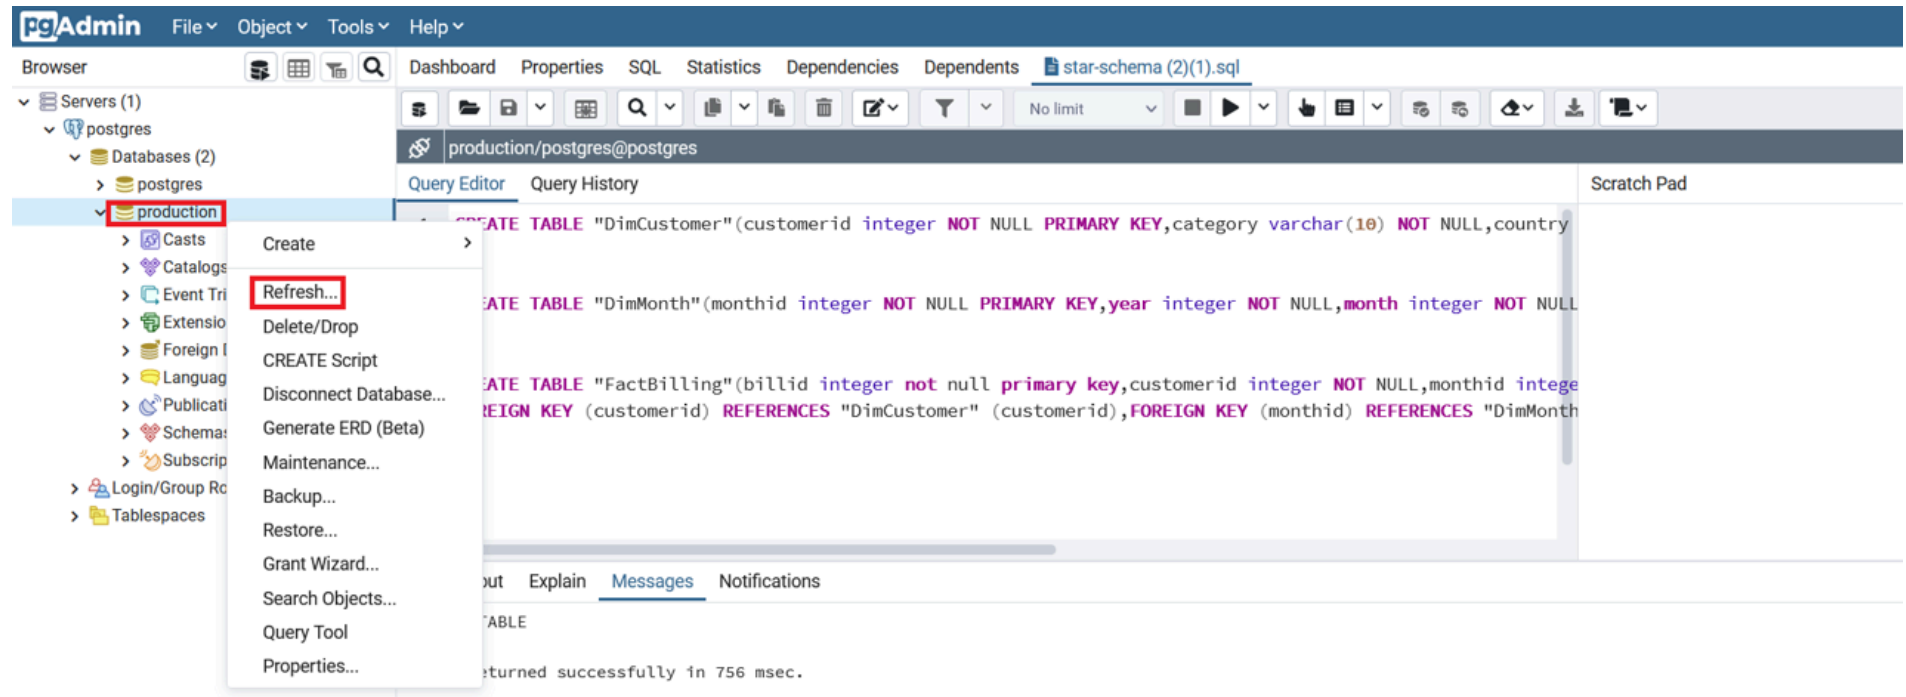
3. Once you click on the **X** icon a new page appears with the file **star-schema.sql**. Select the **star-schema.sql** file from the list and click the **Select** tab.



4. Once the file opens up click on the **Run** option to execute the **star-schema.sql** file.



5. Next, right-click on the **Production** database and click on the **Refresh** option from the dropdown.



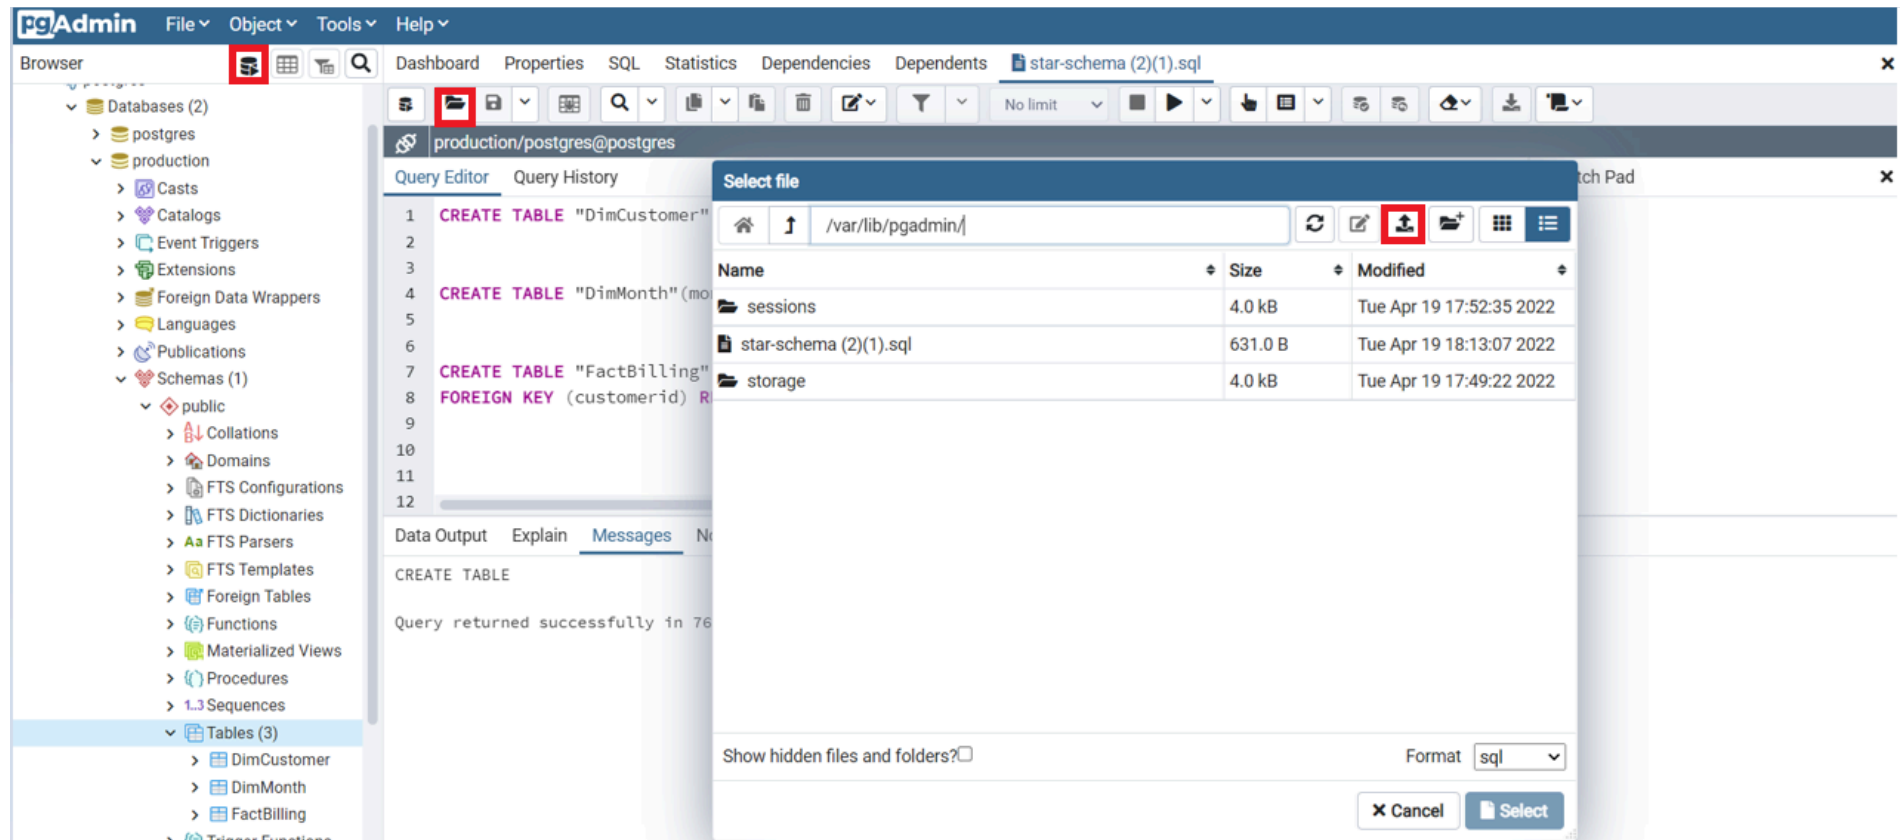
After the database is refreshed the 3 tables (DimCustomer, DimMonth, FactBilling) are created under the **Databases > Production > Schemas > Public > Tables**.

The screenshot shows the pgAdmin interface with the following elements:

- Browser:** A tree view on the left showing the database structure. The 'production' database is selected, and the 'public' schema is expanded. The 'Tables (3)' folder is highlighted, showing 'DimCustomer', 'DimMonth', and 'FactBilling'.
- Query Editor:** A central pane showing a SQL script to create three tables: 'DimCustomer', 'DimMonth', and 'FactBilling'. The script includes primary key and foreign key definitions.
- Messages:** A pane at the bottom showing the execution results of the query, indicating it was successful and returned in 769 msec.

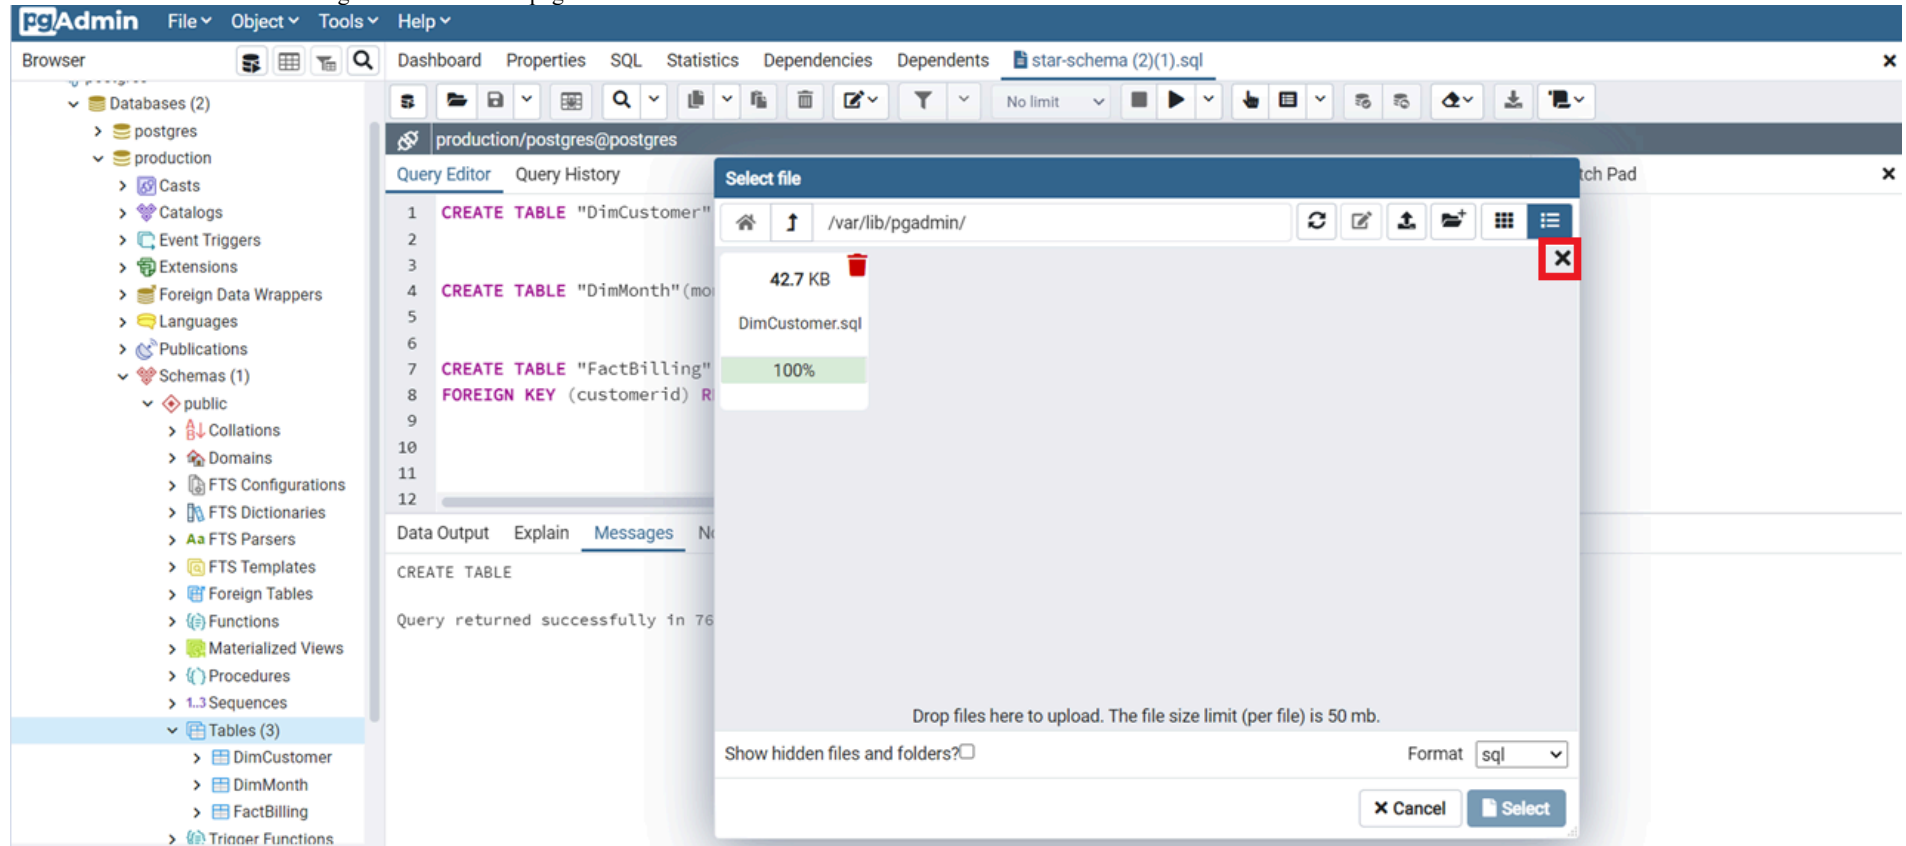
Task C: Load tables

1. Click on **Query** tool and then click **Open** file and click on **Upload** icon.

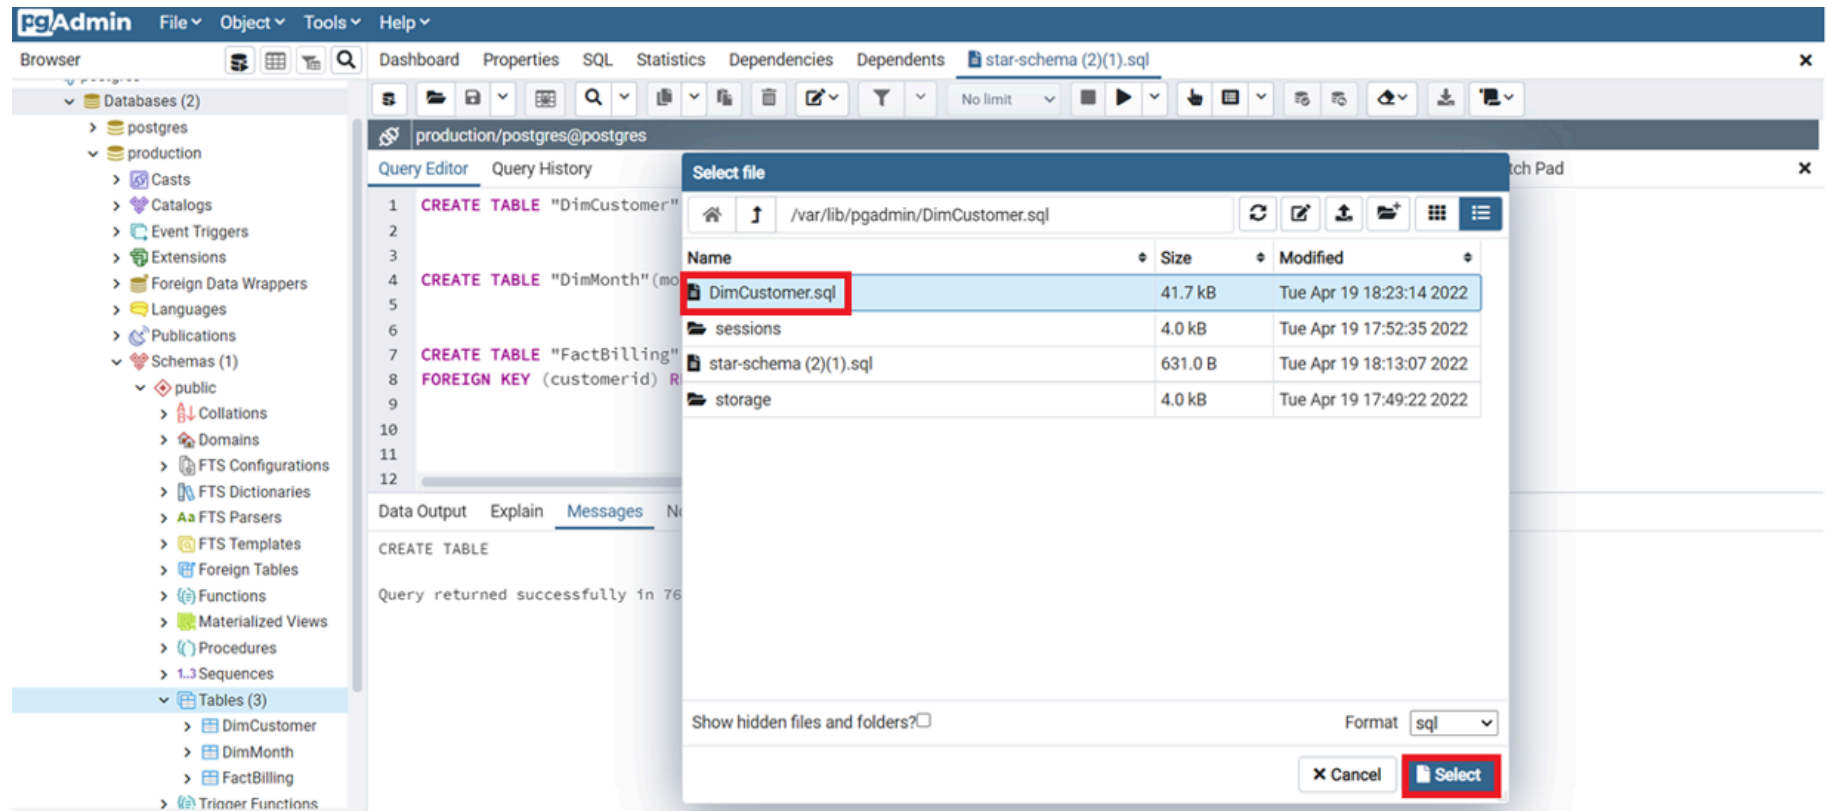


2. In the new blank page that appears drag and drop the **DimCustomer.sql** file inside the blank page. Once the **DimCustomer.sql** file is successfully loaded.

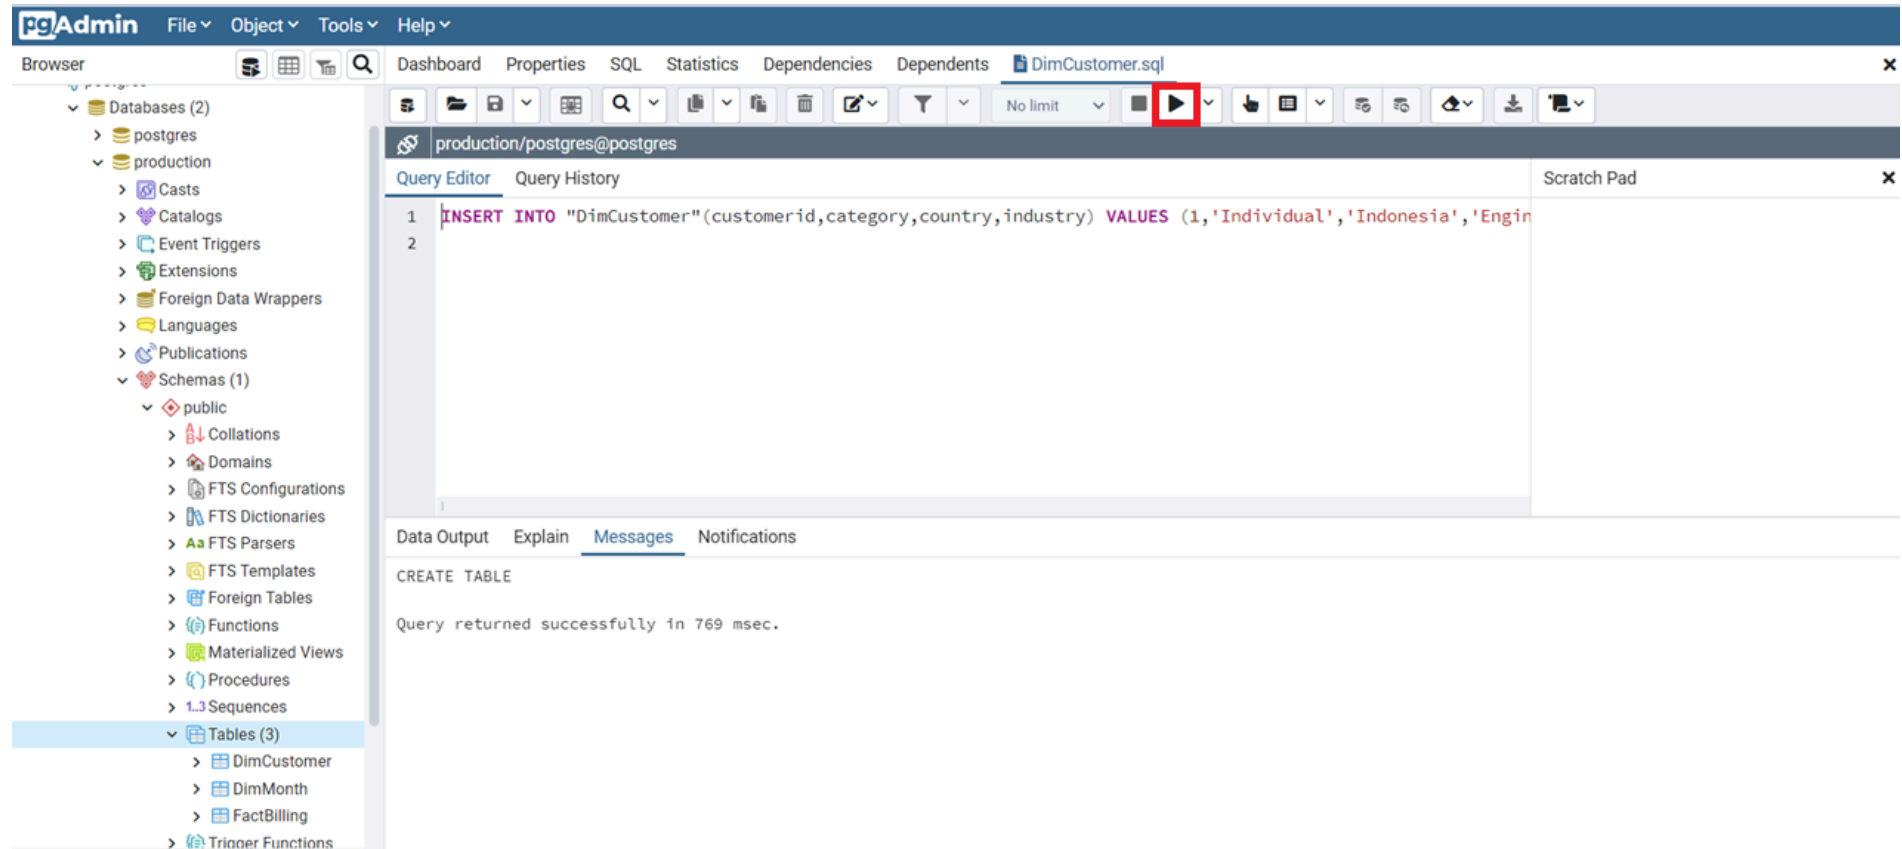
Click on the small **X** icon on the right hand side of the page as shown in the screenshot.



3. Once you click on the **X** icon a new page appears with the file **DimCustomer.sql**. Select the **DimCustomer.sql** file from the list and click on **Select** tab.



4. Once the file opens up, click on the **Run** option to execute the **DimCustomer.sql** file.



Note: Repeat the steps as given in Task C to upload the remaining sql files to insert data in **DimMonth** and **FactBilling**.

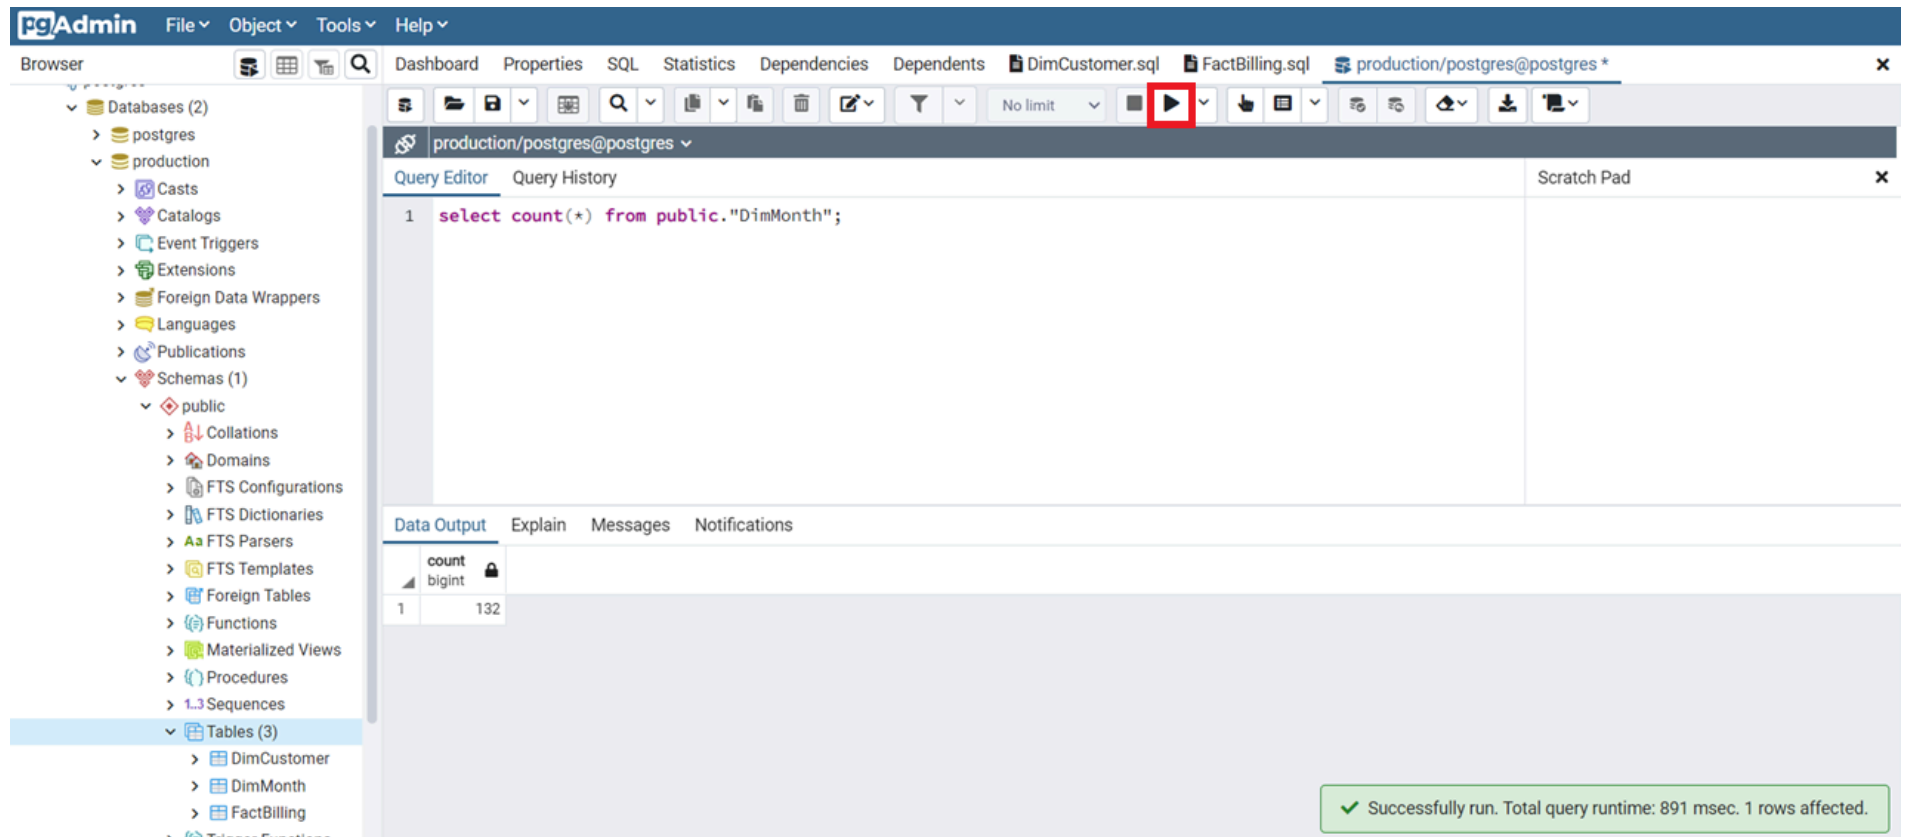
5. Let's run the command below on the PostgreSQL Tool.

1. 1

1. select count(*) from public."DimMonth";

Copied!

You should see an output as seen in the image below.



The screenshot shows the pgAdmin 4 web interface. On the left is a tree view of the database structure, including 'production/postgres@postgres' and 'public' schema. The main area is the 'Query Editor' with a SQL query: `select count(*) from public.\"DimMonth\";`. Below the editor is the 'Data Output' tab showing a single row with the count 132. A green notification box at the bottom right states: 'Successfully run. Total query runtime: 891 msec. 1 rows affected.'

You are encouraged to run more SQL queries.

Practice exercises

Problem 1: Using the PostgreSQL tool, find the count of rows in the table FactBilling

- Click here for Hint
- ▼ Click here for Solution

1. 1

1. `select count(*) from public.\"FactBilling\";`

Copied!

Problem 2: Using the PostgreSQL tool, create a simple Materialized views named avg_customer_bill with fields customerid and averagebillamount.

- Click here for Hint
- ▼ Click here for Solution

1. 1
2. 2
3. 3

```
4. 4
5. 5

1. CREATE MATERIALIZED VIEW avg_customer_bill (customerid, averagebillamount) AS
2. (select customerid, avg(billedamount)
3. from public."FactBilling"
4. group by customerid
5. );
```

Copied!

Click the **Run All** Button to run the statement. You should see status as **Success** in the **Result** section.

Problem 3: Refresh the newly created Materialized views

- Click here for Hint
- ▼ Click here for Solution

```
1. 1

1. REFRESH MATERIALIZED VIEW avg_customer_bill;
```

Copied!

Problem 4: Using the newly created Materialized views find the customers whose average billing is more than 11000.

- Click here for Hint
- ▼ Click here for Solution

```
1. 1

1. select * from avg_customer_bill where averagebillamount > 11000;
```

Copied!

Congratulations! You have successfully finished the Populating a Data Warehouse lab.

Author

Amrutha Rao

© IBM Corporation 2022. All rights reserved.